# Linear models

## How do they work?

Samuel Robinson, Ph.D.

October 8, 2020

# Motivation

- *I have some bivariate data (2 things measured per row), and I want to know if they're related to each other*
- *I have 2+ groups of data, and I want to know whether the means are different*

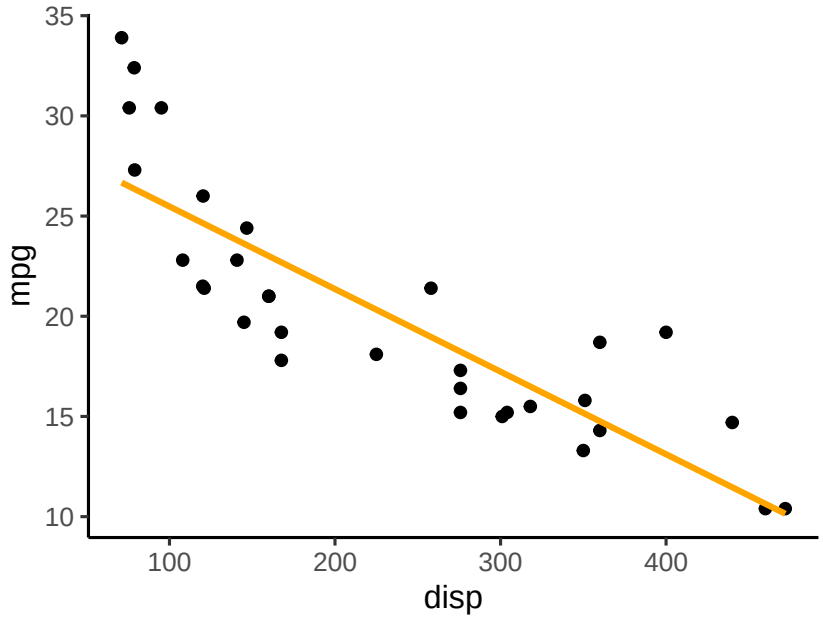


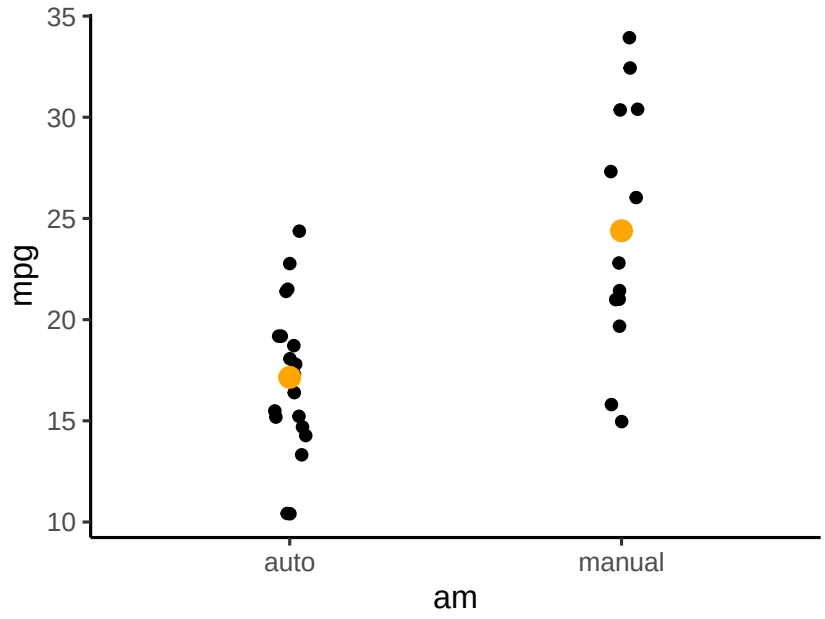

# Model terminology

- All linear models take the form:

$$\hat{y} = b_0 + b_1 x_1 + b_2 x_2 ... + b_i x_i$$
$$y \sim Normal(\hat{y}, \sigma)$$

- $y$ is the thing you're interested in predicting
- $\hat{y}$ is the *predicted value* of $y$
- $x_1 ... x_i$ are *predictors* of $y$
- $b_1 ... b_i$ are *coefficients* for each predictor $x_i$
- $b_0$ is the *intercept*, a coefficient that doesn't depend on predictors
- $y \sim Normal(\hat{y}, \sigma)$ means:
  - "$y$ follows a Normal distribution with mean $\hat{y}$ and SD $\sigma$"

This may look terrifying, but let's use a simple example:

# Example

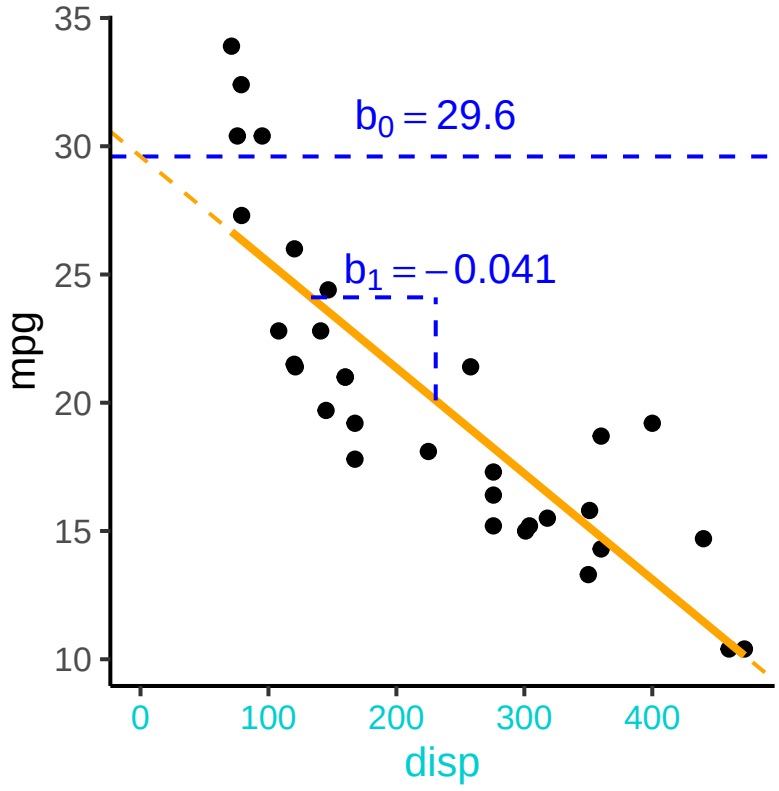


$$\hat{mpg} = b_0 + b_1 disp$$

$$mpg \sim Normal(\hat{mpg}, \sigma)$$

- *mpg* is the thing you're interested in predicting
- $\hat{mpg}$ is the *predicted value* of *mpg*
- *disp* is the *predictor* of *mpg*
- $b_0$ is the *intercept*, $b_1$ is the *coefficient* for *disp*
- $mpg \sim Normal(\hat{mpg}, \sigma)$ means:
  - "*mpg* follows a Normal distribution with mean $\hat{mpg}$ and SD $\sigma$"
- $\sigma$ isn't displayed on the figure. Where is it?

# Example (cont.)

$\sigma$ isn't displayed on the figure. Where is it?

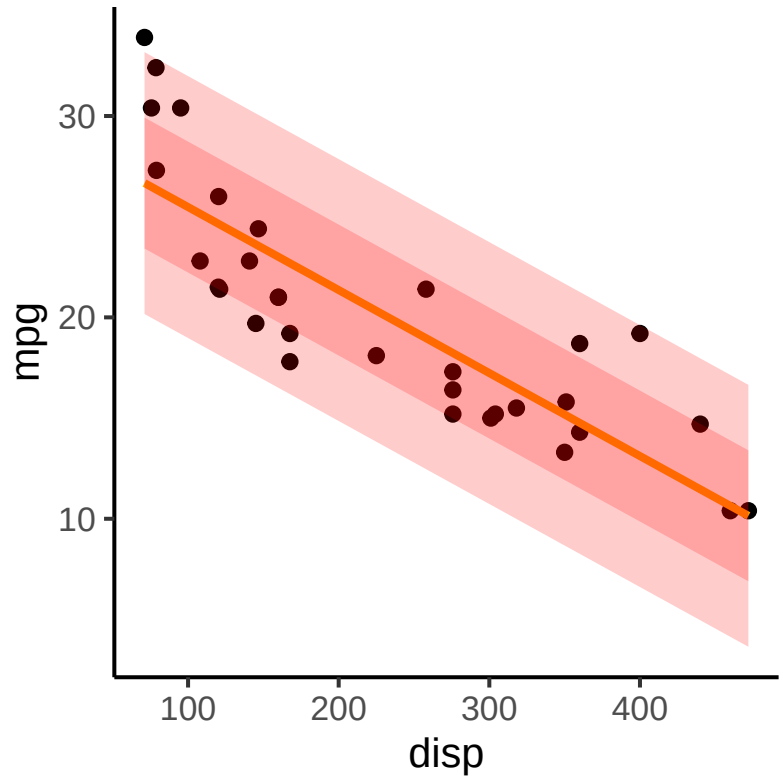


- $\sigma$ is the "leftover" or "residual" variance
- i.e. variation between samples that the model couldn't explain
- Since $y \sim Normal(\hat{y}, \sigma)$, this means that points are normally distributed around the *entire line* of $\hat{y}$

# How do I get R to fit this model?

`lm` is one of the main functions used for linear modeling:

```
#Formula structure: y ~ x
mod1 <- lm(mpg ~ disp, #mpg depends on disp
           data = mtcars) #Name of the dataframe containing mpg & disp
summary(mod1)
```

```
##
## Call:
## lm(formula = mpg ~ disp, data = mtcars)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -4.8922 -2.2022 -0.9631  1.6272  7.2305
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 29.599855   1.229720  24.070  < 2e-16 ***
## disp        -0.041215   0.004712  -8.747 9.38e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.251 on 30 degrees of freedom
## Multiple R-squared:  0.7183, Adjusted R-squared:  0.709
## F-statistic: 76.51 on 1 and 30 DF,  p-value: 9.38e-10
```

For a detailed breakdown of `lm`'s output, click here

# Simulate data

Now that we know how linear models work, we can simulate our own data:

```r
#Parameters:

b0 <- 1 #Intercept
b1 <- 2 #Slope
sigma <- 3 #SD

#Make up some data:

x <- 0:30 #Predictor values

#Predicted y values
pred_y <- b0 + b1*x

#Add "noise" around pred_y
actual_y <- rnorm(n = length(pred_y),
                  mean = pred_y,
                  sd= sigma)
```

```r
#Plot the data we just made
plot(x,pred_y,col='orange',pch=19,
     ylab='y')
points(x,actual_y,col='black',pch=19)
```

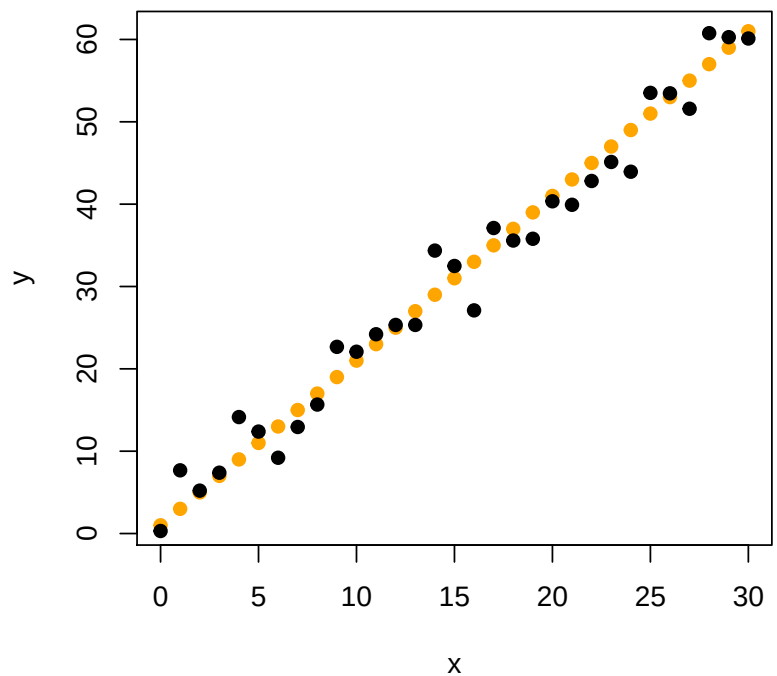

# Fit a model from simulated data

How does R do at finding the coefficients?

Remember: $b_0 = 1, b_1 = 2, \sigma = 3$

```
#Put the simulated data into a dataframe
fakeDat <- data.frame(x = x, y = actual_y, pred = pred_y)
mod1sim <- lm(y ~ x, data = fakeDat) #Fit a linear model
summary(mod1sim)
```

```
##
## Call:
## lm(formula = y ~ x, data = fakeDat)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -5.7568 -1.7623 -0.2176  1.9419  5.3572
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.02974    1.00445   2.021   0.0526 .
## x            1.92670    0.05751  33.499   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.864 on 29 degrees of freedom
## Multiple R-squared:  0.9748, Adjusted R-squared:  0.9739
## F-statistic:  1122 on 1 and 29 DF,  p-value: < 2.2e-16
```

# What about categorical data?

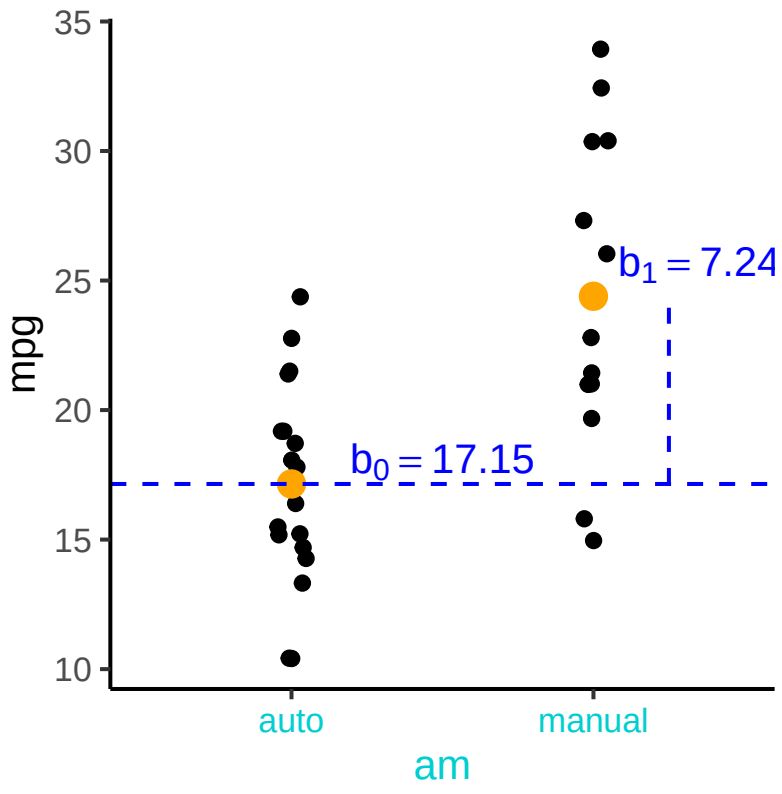


This uses *exactly the same* math!

- *mpg* is the thing you're interested in predicting
- $\hat{mpg}$ is the *predicted value* of *mpg*
- *am* is the *predictor* of *mpg*
  - set of 0s and 1s, not continuous
- $b_0$ is the *intercept*, $b_1$ is the *coefficient* for *am*
- Where is $\sigma$?

$$\hat{mpg} = b_0 + b_1 am$$

$$mpg \sim Normal(\hat{mpg}, \sigma)$$

# How do I get R to fit this model?

Syntax is exactly the same for this model

```r
#Formula structure: y ~ x
mod2 <- lm(mpg ~ am, #mpg depends on am
           data = mtcars) #Name of the dataframe containing mpg & am
summary(mod2)
```

```
##
## Call:
## lm(formula = mpg ~ am, data = mtcars)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -9.3923 -3.0923 -0.2974  3.2439  9.5077
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)   17.147      1.125  15.247 1.13e-15 ***
## am             7.245      1.764   4.106 0.000285 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.902 on 30 degrees of freedom
## Multiple R-squared:  0.3598, Adjusted R-squared:  0.3385
## F-statistic: 16.86 on 1 and 30 DF,  p-value: 0.000285
```

# A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:
  - ~~Steal~~Borrow my code, and change the predictor from continuous to discrete
  - Useful command: `rep` (replicate)
    - e.g. `rep(x=c(0,1),each=10)`
  - Useful command: `rnorm` (generate normally-distributed data)
    - e.g. `rnorm(n=100,mean=0,sd=1)`
- Use `lm` to fit a model to the data you just simulated
  - How does R do at guessing your coefficients?