

Mixed effects models

Wheels within wheels

Samuel Robinson, Ph.D.

December 10, 2020

Motivation

- What are mixed effects models?
 - Scary math (matrix algebra)
 - Variance partitioning
 - Fixed effects vs. Random effects
- Working with random effects
 - What should be a random effect?
 - Model validation
- Exercise

What are mixed effects models?

Many different names:

- ① Mixed effects models
- ② Random effects models
- ③ Hierarchical models
- ④ Empirical/Bayesian hierarchical models
- ⑤ Latent variable models
- ⑥ Split-plot models¹

I like the term *heirarchical models*, as this is the closest to what I will teach you

¹Earlier form of variance partitioning

Scary math

Unfortunately, we need a review of matrix algebra in order to explain this:

- This is a matrix:

$$A = \begin{bmatrix} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 9 \end{bmatrix}$$

- This is a vector

$$b = \begin{bmatrix} 1 & 2 & 3 \end{bmatrix}$$

- Multiplying them looks like this:

$$A \times b = Ab = 1 \times \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} + 2 \times \begin{bmatrix} 4 \\ 5 \\ 6 \end{bmatrix} + 3 \times \begin{bmatrix} 7 \\ 8 \\ 9 \end{bmatrix} = \begin{bmatrix} 30 \\ 36 \\ 42 \end{bmatrix}$$

Why do we call them “linear models”?

Involves a *linear mapping* of coefficients onto a model matrix

Coefficients:

$$\beta = \begin{bmatrix} 0.1 & 1.8 & -0.03 \end{bmatrix}$$

Model matrix:

$$X = \begin{bmatrix} 1 & 1 & 10 \\ 1 & 1 & 12 \\ 1 & 0 & 9 \\ \vdots & \vdots & \vdots \end{bmatrix}$$

Multiplying them looks like:

$$\hat{y} = X\beta = \begin{bmatrix} 1.60 \\ 1.54 \\ -0.17 \\ \vdots \end{bmatrix}$$

This is exactly what R does to fit models:

```
head(dat)
```

```
##           y           x site
## 1 -2.192241 -4.248450    c
## 2  5.735165  5.766103    n
## 3  3.834614 -1.820462    o
## 4 10.403239  7.660348    h
## 5 15.411355  8.809346    a
## 6 -9.085650 -9.088870    g
```

```
m1 <- lm(y~x,data=dat) #Use variable x to predict y
summary(m1)
```

```
##
## Call:
## lm(formula = y ~ x, data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -25.2451  -4.4342   0.5114   6.2092  18.1923
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.0753    0.6530   1.647   0.102
## x             0.7394    0.1141   6.482 1.27e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.996 on 148 degrees of freedom
## Multiple R-squared:  0.2211, Adjusted R-squared:  0.2159
## F-statistic: 42.02 on 1 and 148 DF,  p-value: 1.267e-09
```

This is exactly what R does to fit models (cont.):

```
head(model.matrix(m1))
```

```
##      (Intercept)          x
## 1             1 -4.248450
## 2             1  5.766103
## 3             1 -1.820462
## 4             1  7.660348
## 5             1  8.809346
## 6             1 -9.088870
```

```
coef(m1)
```

```
## (Intercept)          x
## 1.0753139    0.7393775
```

```
pred2 <- model.matrix(m1) %*% coef(m1) #predicted = matrix * coeffs
head(data.frame(pred1=predict(m1),pred2)) #same thing!
```

```
##      pred1      pred2
## 1 -2.0658941 -2.0658941
## 2  5.3386403  5.3386403
## 3 -0.2706944 -0.2706944
## 4  6.7392027  6.7392027
## 5  7.5887456  7.5887456
## 6 -5.6447919 -5.6447919
```

Groups are coded by “dummy variables” (0s and 1s)

```
m2 <- lm(y~site,data=dat) #Use variable site to predict y  
head(model.matrix(m2)) #0s and 1s used to identify groups
```

```
## (Intercept) siteb sitec sited sitee sitef siteg siteh sitei sitej sitek sitel  
## 1 1 0 1 0 0 0 0 0 0 0 0  
## 2 1 0 0 0 0 0 0 0 0 0 0  
## 3 1 0 0 0 0 0 0 0 0 0 0  
## 4 1 0 0 0 0 0 0 1 0 0 0  
## 5 1 0 0 0 0 0 0 0 0 0 0  
## 6 1 0 0 0 0 0 1 0 0 0 0  
## sitem siten siteo  
## 1 0 0 0  
## 2 0 1 0  
## 3 0 0 1  
## 4 0 0 0  
## 5 0 0 0  
## 6 0 0 0
```

```
coef(m2) #This uses the 1st site as the "control" group
```

```
## (Intercept)      siteb      sitec      sited      sitee      sitef  
## 0.1493889 -0.3476999 -3.1293722  4.7334445 -2.5721233 14.7029246  
##      siteg      siteh      sitei      sitej      sitek      sitel  
## 1.7083447  2.3612282 -7.8406290 -6.4931198  4.7413122 -3.9443779  
##      sitem      siten      siteo  
## 5.8387513  0.8629803  1.9148871
```


Structure of LMs. . . now with matrices!

- All linear models take the form:

$$\hat{y} = X\beta = b_0\mathbf{1} + b_1x_1\ldots + b_ix_i$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is a vector of data you want to predict
- \hat{y} is a vector of *predicted values* for y
- $X = \{\mathbf{1}, x_1\ldots\}$ is a matrix of *predictors* for y
- $\beta = \{b_0, b_1, \ldots\}$ is a vector of *coefficients*
- $y \sim \text{Normal}(\hat{y}, \sigma)$ means:
 - “ y follows a Normal distribution with mean \hat{y} and SD σ ”

Fixed effects vs. Random effects

Say that X is a model matrix coding for 10 sites², and y is something we're interested in predicting

$$\hat{y} = b_0 + X\beta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- Site coefficients (β) are unrelated to each other
- σ is the SD of *residuals*
- Site is a **fixed effect**

$$\hat{y} = b_0 + X\zeta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$
$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Site coefficients (β) are related to each other via a *Normal* distribution
- σ is the SD of *residuals*, σ_{site} is the SD of *sites*
- Site is a **random effect**

²Intercept is a separate variable

Mixed effects = fixed + random effects

A mixed effects model has both **fixed** and **random** effects

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients
- U = random effects matrix (e.g. sites)
- ζ = random effects coefficients
- $\sigma, \sigma_{\text{site}}$ = variance terms

Mixed effect model example

Let's go back to our earlier example:

- We're interested in predicting y using x (fixed effects)
- Data was collected at a number of *sites*, which may affect y "somehow"
- Effect of each site is normally distributed

```
head(dat)
```

```
##           y           x site
## 1 -2.192241 -4.248450    c
## 2  5.735165  5.766103    n
## 3  3.834614 -1.820462    o
## 4 10.403239  7.660348    h
## 5 15.411355  8.809346    a
## 6 -9.085650 -9.088870    g
```

```
library(lme4) #Mixed effects library
mm1 <- lmer(y ~ x + (1|site),data=dat) #site is fit as "random intercepts"
```

Mixed effect model example

```
summary(mm1)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + (1 | site)
## Data: dat
##
## REML criterion at convergence: 1003.2
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.14013 -0.55056  0.09106  0.53825  2.38279
##
## Random effects:
## Groups   Name      Variance Std.Dev.
## site    (Intercept) 23.72    4.870
## Residual                39.38    6.276
## Number of obs: 150, groups: site, 15
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  0.99485    1.37791    0.722
## x            0.71548    0.09409    7.604
##
## Correlation of Fixed Effects:
## (Intr)
## x -0.006
```


Results from lmer model:

- Random effects:
 - *residual* and *site* variance (σ , σ_{site})
- Fixed effects:
 - Intercept and slope estimates (β)
 - No d.f. and p-value³

³'lme4' author doesn't think they can be calculated. I largely agree

Mixed effect model results



- In a *random intercepts* model, the regression line of x on y is allowed to move up or down around the main regression line for each site
- These changes in intercepts are *normally distributed*



Mixed effect model results (cont.)

For plotting, we want a partial effects plot that marginalizes across sites (i.e. “What does the trend look like at the average site?”)

- `ggpredict` works well for this. If you want partial residuals, you'll have to add them in yourself using `predict` and `residual`



Random slope + intercept

Suppose that y wasn't just higher or lower at each site, but that the effect of x on y was higher or lower at each site

$$\hat{y} = X\beta + U\zeta_{int} + x_i U\zeta_{slope}$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta_{int} \sim \text{Normal}(0, \sigma_{int})$$

$$\zeta_{slope} \sim \text{Normal}(0, \sigma_{slope})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- x_i = individual fixed effect (e.g. temperature)
- β = fixed effects coefficients
- U = random effects matrix (e.g. sites)
- $\zeta_{int}, \zeta_{slope}$ = random intercept and slope coefficients
- $\sigma, \sigma_{int}, \sigma_{slope}$ = variance terms

Random slope and intercept example:

```
#Intercept varies with site, and slope of x can  
# also vary with site (both hierarchical)  
mm2 <- lmer(y ~ x + (x|site),data=dat)  
summary(mm2)
```


```
## Linear mixed model fit by REML ['lmerMod']  
## Formula: y ~ x + (x | site)  
## Data: dat  
##  
## REML criterion at convergence: 868.9  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.7003 -0.6290  0.0972  0.5622  2.4476   
##  
## Random effects:  
## Groups   Name                Variance Std.Dev. Corr   
## site    (Intercept)  30.695    5.540               
##         x              1.012    1.006    0.61        
## Residual              10.781    3.283               
## Number of obs: 150, groups: site, 15  
##  
## Fixed effects:  
##              Estimate Std. Error t value        
## (Intercept)   0.2376     1.4673    0.162        
## x              0.4858     0.2676    1.815        
##  
## Correlation of Fixed Effects:  
##      (Intr)   
## x  0.576
```

Results from lmer model:

- Random effects:
 - *residual, slope, and site* variance (σ , σ_{int} , σ_{slope})
 - Correlation b/w intercept and slope = 0.61
 - Sites with higher intercept *also* have a higher slope
- Fixed effects:
 - Intercept and slope estimates

Mixed effect model results

- Regression line of x on y is allowed to move up or down around the main regression line for each site (random intercepts)
- Slope of regression line can be more or less steep for each site (random slopes)
- Changes in intercepts and slopes are *normally distributed*, and in this example, are correlated with each other



Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can't I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!
 - You lose 1 d.f. for each site = Type II error \uparrow = You may not find the fixed effect of interest, even if it's there!
 - Sites with low sample sizes may cause your models to break
 - People⁴ who have read statistics books published after 1980 may ask questions
- However, if you have a low number of sites (1-10), fixed effects may work better
 - Hard to estimate σ_{site} if number of sites is low
 - Partial pooling doesn't really help if no other sites are available to “borrow strength” from
 - Easier to interpret (p-values, ANOVA, etc.)

⁴e.g. me