

Mixed effects models 2

Even more bells and whistles

Samuel Robinson, Ph.D.

December 17, 2020

Motivation


- How do I check if model results are valid?
 - Residual checks
 - Hypothesis testing
- What if my response variable is non-normal?
 - Generalized linear mixed models (GLMMs)
- Sampling over time or space
 - “Continuous” random effects
- Christmas-themed exercise!

Mixed effect model example

Let's go back to our earlier example:






```
lmer(y ~ x + (x|site), data = dat)
```

- We're interested in predicting y using x (fixed effects)
- Data was collected at a number of *sites*, which may affect y
- Effect of each site is normally distributed (random intercept)
- Effect of site on slope of x is normally distributed (random slope)



Validation

- Similar to linear models, but we *also* check whether the random intercepts are normally distributed



Hypothesis testing

Is this fixed effect important? (e.g. ANOVA)

- Use likelihood-based test via drop1 (likelihood ratio test, AIC)
- Be careful to fit model with REML = FALSE!

```
lmm1 <- update(lmm1,REML=FALSE) #Refit model using ML rather than REML  
drop1(lmm1,test='Chisq') #x has a strong effect
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## y ~ x + (x | site)
```

```
##          npar      AIC      LRT  Pr(Chi)
```

```
## <none>      913.08
```

```
## x          1 917.85 6.7697 0.009272 **
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Hypothesis testing (cont.)

How do I know this effect is different from x ? - Use Wald Z-test (2-sided p-value from Z-test)

```
##                x
## 0.004659069
```

- `glht` from `library(multcomp)` works with `lmer` models if you are comparing between coefficients (e.g. “Is treatment A different from B and C?”)

What if my response variable is non-normal?

- Linear model (LM)

$$\hat{y} = X\beta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- Linear mixed effects model (LMM)

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Generalized linear model (GLM)

$$\text{logit}(\hat{\phi}) = X\beta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

- Generalized linear mixed effects model (GLMM)

$$\text{logit}(\hat{\phi}) = X\beta + U\zeta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

How do I fit GLMMs?

- `glmer` and `glmer.nb` from `library(lme4)` work for Binomial, Poisson, and Negative Binomial data

```
library(lme4)
glmm1 <- glmer.nb(y2~x+(x|site),data=dat) #Negative binomial GLMM
summary(glmm1) #glmer.nb takes a LONG time to run
```

- `glmmTMB` from `library(glmmTMB)` works for those above, *plus* a bunch of others
 - e.g. Zero-inflation, Beta-binomial, Spatial Models

```
library(glmmTMB)
glmm2 <- glmmTMB(y2~x+(x|site),data=dat,family=nbinom2())
summary(glmm2) #Similar results, but quicker
```


Fitting GLMMs - glmer.nb

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: Negative Binomial(5.1295) ( log )
## Formula: y2 ~ x + (x | site)
## Data: dat
##
##          AIC          BIC    logLik deviance df.resid
##      627.8      646.3   -307.9   615.8      154
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.3745 -0.7098 -0.3946  0.5108  2.5367
##
## Random effects:
##  Groups Name      Variance Std.Dev. Corr
##  site (Intercept) 1.43502  1.1979
##      x           0.02878  0.1697   0.92
## Number of obs: 160, groups:  site, 16
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.32745    0.32166   1.018   0.3087
## x           0.10830    0.04681   2.314   0.0207 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##  (Intr)
## x 0.799
```

Fitting GLMMs - glmmTMB

```
## Family: nbinom2 ( log )
## Formula:          y2 ~ x + (x | site)
## Data: dat
##
##      AIC      BIC   logLik deviance df.resid
##    627.8    646.2   -307.9    615.8      154
##
## Random effects:
##
## Conditional model:
##   Groups Name      Variance Std.Dev. Corr
##   site  (Intercept) 1.43543  1.1981
##         x           0.02892  0.1701   0.92
## Number of obs: 160, groups: site, 16
##
## Overdispersion parameter for nbinom2 family (): 5.12
##
## Conditional model:
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.34132    0.32172   1.061   0.2887
## x           0.11026    0.04697   2.348   0.0189 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual checks on glmmTMB objects

- Extract residuals and make your own plots, or use `simulateResiduals` from `library(DHARMA)` (see here)
- DHARMA also has useful functions for checking overdispersion and zero-inflation (found here)

```
## This is DHARMA 0.3.3.0. For overview type '?DHARMA'. For recent changes, type news(package = 'DHARMA')
```

DHARMA residual diagnostics

